

Lab 7: Relational Data

Sunghwu Song

March 19, 2021

Introduction

The main purpose of this lab is to practice data join skills from Chapter 10. The functions and their purposes are listed as follows:

- `inner_join()` Keeps observations appear in both datasets.
- `left_join()` Keeps all observations in left dataset.
- `right_join()` Keeps all observations in right dataset.
- `full_join()` Keeps all observations in both datasets.
- `semi_join()` Keeps all observations in left dataset that have a match in right dataset.
- `anti_join()` Drops all observations in left dataset that have a match in right dataset.

You will need to modify the code chunks so that the code works within each of chunk (usually this means modifying anything in ALL CAPS). You will also need to modify the code outside the code chunk. When you get the desired result for each step, change `Eval=F` to `Eval=T` and knit the document to HTML to make sure it works. After you complete the lab, you should submit your HTML file of what you have completed to Sakai before the deadline.

Excercises

Part 1

In part 1, you will practice the skills using the datasets from the R package **Lahman**. This database includes data related to baseball teams. It includes summary statistics about how the players performed on offense and defense for several years. It also includes personal information about the players.

The **Batting** data frame contains the offensive statistics for all players for many years. You can see, for example, the top 10 hitters in 2016 by running this code: (For more details of the dataset run `?Batting` in console.)

```
top <- Batting %>%
  filter(yearID == 2016) %>%
  arrange(desc(HR)) %>%
  slice(1:10)

top
```

```
##      playerID yearID stint teamID lgID   G  AB   R   H  X2B  X3B  HR  RBI  SB  CS  BB
## 1  trumbma01   2016     1   BAL   AL 159 613  94 157  27   1 47 108  2  0 51
## 2   cruzne02   2016     1   SEA   AL 155 589  96 169  27   1 43 105  2  0 62
## 3  daviskh01   2016     1   OAK   AL 150 555  85 137  24   2 42 102  1  2 42
## 4  doziebr01   2016     1   MIN   AL 155 615 104 165  35   5 42  99 18  2 61
## 5  encared01   2016     1   TOR   AL 160 601  99 158  34   0 42 127  2  0 87
## 6  arenano01   2016     1   COL   NL 160 618 116 182  35   6 41 133  2  3 68
## 7  cartech02   2016     1   MIL   NL 160 549  84 122  27   1 41  94  3  1 76
## 8  frazito01   2016     1   CHA   AL 158 590  89 133  21   0 40  98 15  5 64
## 9  bryankr01   2016     1   CHN   NL 155 603 121 176  35   3 39 102  8  5 75
## 10 canoro01   2016     1   SEA   AL 161 655 107 195  33   2 39 103  0  1 47
##      SO  IBB  HBP  SH  SF  GDP
## 1  170   1   3   0   0   14
## 2  159   5   9   0   7   15
## 3  166   0   8   0   5   19
## 4  138   6   8   2   5   12
## 5  138   3   5   0   8   22
## 6  103  10   2   0   8   17
## 7  206   1   9   0  10   18
## 8  163   1   4   1   7   11
## 9  154   5  18   0   3    3
## 10 100   8   8   0   5   18
```

But who are these players? We see an ID, but not the names. The player names are in this table

```
head(Master,5)
```

```
##      playerID birthYear birthMonth birthDay birthCountry birthState  birthCity
## 1  aardsda01     1981         12        27          USA          CO      Denver
## 2  aaronha01     1934          2         5          USA          AL      Mobile
## 3  aaronto01     1939          8         5          USA          AL      Mobile
## 4  aasedo01     1954          9         8          USA          CA      Orange
## 5  abadan01     1972          8        25          USA          FL  Palm Beach
##      deathYear deathMonth deathDay deathCountry deathState deathCity nameFirst
## 1         NA         NA         NA          <NA>          <NA>          <NA>    David
## 2         NA         NA         NA          <NA>          <NA>          <NA>    Hank
## 3     1984          8        16          USA          GA    Atlanta    Tommie
## 4         NA         NA         NA          <NA>          <NA>          <NA>      Don
## 5         NA         NA         NA          <NA>          <NA>          <NA>     Andy
##      nameLast      nameGiven weight height bats  throws      debut  finalGame
## 1  Aardsma    David Allan   215    75   R      R 2004-04-06 2015-08-23
## 2   Aaron    Henry Louis   180    72   R      R 1954-04-13 1976-10-03
## 3   Aaron    Tommie Lee   190    75   R      R 1962-04-10 1971-09-26
## 4   Aase Donald William   190    75   R      R 1977-07-26 1990-10-03
## 5   Abad  Fausto Andres   184    73   L      L 2001-09-10 2006-04-13
##      retroID  bbrefID  deathDate  birthDate
## 1  aardd001  aardsda01          <NA> 1981-12-27
## 2  aaroh101  aaronha01          <NA> 1934-02-05
## 3  aarot101  aaronto01 1984-08-16 1939-08-05
## 4  aased001  aasedo01          <NA> 1954-09-08
## 5  abada001  abadan01          <NA> 1972-08-25
```

We can see column names `nameFirst` and `nameLast` in table `Master`.

Question 1

1. Use the `left_join` function to create a data frame called `top1`, which contains information of the 10 top home run hitters. The table should have the following columns: `playerID`, `nameFirst`, `nameLast`, and number of home runs (`HR`). (1 Point)

```
top1 = top %>%
  left_join(Master, by = c("playerID" = "playerID")) %>%
  select('playerID', 'nameFirst', 'nameLast', 'HR')

top1
```

```
##   playerID nameFirst nameLast HR
## 1 trumbma01      Mark    Trumbo 47
## 2 cruzne02    Nelson      Cruz 43
## 3 daviskh01    Khris     Davis 42
## 4 doziebr01    Brian     Dozier 42
## 5 encared01    Edwin Encarnacion 42
## 6 arenano01    Nolan     Arenado 41
## 7 cartech02    Chris     Carter 41
## 8 frazito01    Todd      Frazier 40
## 9 bryankr01    Kris      Bryant 39
## 10 canoro01 Robinson    Cano 39
```

Question 2

Data `Salaries` contains the baseball player salary data.

```
head(Salaries,5)
```

```
##   yearID teamID lgID  playerID salary
## 1  1985    ATL   NL  barkele01 870000
## 2  1985    ATL   NL  bedrost01 550000
## 3  1985    ATL   NL  benedbr01 545000
## 4  1985    ATL   NL  campri01 633333
## 5  1985    ATL   NL  ceronri01 625000
```

2. You may be curious about the salaries of the top 10 hitters in 2016 (4 Points):
 - Now create a new data frame called `top2` by adding top 10 hitters' salaries to `top1` and including only `nameFirst`, `nameLast`, `teamID`, `HR`, and `salary` columns.
 - Rename the columns to `FirstName`, `LastName`, `Team`, `Homeruns` and `Salary` respectively.
 - Arrange the data frame by `Salary` in descending order.

Note that salaries are different every year so make sure to filter for the year 2016. This time, please use `right_join` to complete the exercise.

```
top2 = Salaries %>%
  filter(yearID == 2016) %>%
  right_join(top1, by="playerID") %>%
  select('FirstName' = 'nameFirst', "LastName" = "nameLast", "Team" = "teamID", "Homeruns" = "HR", "Salary" = "salary")
```

```
arrange(desc(Salary))

top2
```

##	FirstName	LastName	Team	Homeruns	Salary
## 1	Robinson	Cano	SEA	39	24000000
## 2	Nelson	Cruz	SEA	43	14250000
## 3	Edwin	Encarnacion	TOR	42	10000000
## 4	Mark	Trumbo	BAL	47	9150000
## 5	Todd	Frazier	CHA	40	8250000
## 6	Nolan	Arenado	COL	41	5000000
## 7	Brian	Dozier	MIN	42	3000000
## 8	Chris	Carter	MIL	41	2500000
## 9	Kris	Bryant	CHN	39	652000
## 10	Khris	Davis	OAK	42	524500

Part 2

In this part, we will explore relational data from `nycflights13`, which contains four data frames related to the `flights` table that you used in previous assignments.

Question 3

Data `airports` gives information about each airport, such as latitude and longitude, identified by the `faa` airport code.

```
head(airports,5)
```

```
## # A tibble: 5 x 8
##   faa   name                lat   lon   alt   tz dst  tzone
##   <chr> <chr>                <dbl> <dbl> <dbl> <dbl> <chr> <chr>
## 1 04G   Lansdowne Airport      41.1 -80.6  1044   -5 A   America/New_Y~
## 2 06A   Moton Field Municipal Airp~ 32.5 -85.7   264   -6 A   America/Chica~
## 3 06C   Schaumburg Regional     42.0 -88.1   801   -6 A   America/Chica~
## 4 06N   Randall Airport        41.4 -74.4   523   -5 A   America/New_Y~
## 5 09J   Jekyll Island Airport   31.1 -81.4    11   -5 A   America/New_Y~
```

- Based on `flights`, compute the average arrival delay by destination (`dest`) and ignore missing values, then join on the `airports` data frame then show the spatial distribution of delays. (3 Points)

```
delay = flights %>%
  group_by(dest) %>%
  summarise(avg_arr_delay=mean(arr_delay, na.rm = TRUE), .groups='drop') %>%
  inner_join(airports,by=c("dest" = "faa")) %>%
  select("avg_arr_delay", "lat", "lon")
```


Question 4

- Draw a scatterplot with dots representing destination locations and colors of dots representing average arrival delay on US map. (4 Points) (For `coord_quickmap` to work, you need to install `maps` packages. If you haven't installed the package before, please run `install.packages('maps')` in Console.)

```

delay %>%
  ggplot(aes(x = lon, y = lat, color=avg_arr_delay)) +
    borders("state") +
    geom_point() +
    coord_quickmap()

```



Question 5

Data `planes` gives information about each plane, identified by its `tailnum`. Note that `year` column in `planes` represents the year a plane was manufactured, which is different from `year` column in `flights`.

```
head(planes,5)
```

```

## # A tibble: 5 x 9
##   tailnum year type      manufacturer model engines seats speed engine
##   <chr>   <int> <chr>      <chr>      <chr>   <int> <int> <int> <chr>
## 1 N10156  2004 Fixed wing mu~ EMBRAER    EMB-1~     2    55    NA Turbo~~
## 2 N102UW  1998 Fixed wing mu~ AIRBUS INDUST~ A320~~     2   182    NA Turbo~~
## 3 N103US  1999 Fixed wing mu~ AIRBUS INDUST~ A320~~     2   182    NA Turbo~~
## 4 N104UW  1999 Fixed wing mu~ AIRBUS INDUST~ A320~~     2   182    NA Turbo~~
## 5 N10575  2002 Fixed wing mu~ EMBRAER    EMB-1~     2    55    NA Turbo~~

```

5. Use the `planes` data to calculate the `age` of planes, assuming current year is 2013. Keep only `tailnum` and `age` in the output table `plane_ages`. (1 Point)

```

plane_ages <-
  planes %>%
    mutate(age = 2013-year) %>%
    select(tailnum, age)

```


Question 6

6. Is there a relationship between the age of a plane and its delays? (4 Points)

- Join the `plane_ages` with `flights`, keeping observations with matches in both datasets.
- Summarize the average departure delay by plane `age` and ignore missing values.
- Draw a scatterplot of plane age vs. average departure delay.

```
plane_ages %>%  
  inner_join(flights, by = 'tailnum') %>%  
  group_by(age) %>%  
  summarise(avg_dep_delay = mean(dep_delay, na.rm = TRUE), .groups='drop') %>%  
  ggplot(aes(x = age, y = avg_dep_delay)) +  
  geom_point()
```

Warning: Removed 1 rows containing missing values (geom_point).



Question 7

7. What does it mean for a flight to have a missing `tailnum`? (1 Point)

```
flights %>%  
  filter(is.na(tailnum))
```

```
## # A tibble: 2,512 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     2     NA             1545           NA       NA             1910
## 2  2013     1     2     NA             1601           NA       NA             1735
## 3  2013     1     3     NA              857           NA       NA             1209
## 4  2013     1     3     NA              645           NA       NA              952
## 5  2013     1     4     NA              845           NA       NA             1015
## 6  2013     1     4     NA             1830           NA       NA             2044
## 7  2013     1     5     NA              840           NA       NA             1001
## 8  2013     1     7     NA              820           NA       NA              958
## 9  2013     1     8     NA             1645           NA       NA             1838
## 10 2013     1     9     NA              755           NA       NA             1012
## # ... with 2,502 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

Answer: Flights with missing tailnum have no airtime, arrival time, departal time, and other variables that relate to its flights being in the air. Thus the flights with missing tailnum were probably cancelled.

Question 8

8. What do the tail numbers that don't have a matching record in planes have in common? (Hint: one variable explains ~90% of the problems. Check the documentation of `planes` for help.) (2 Points)

```
flights %>%
  anti_join(planes, by='tailnum') %>%
  count(carrier) %>%
  arrange(desc(n))
```

```
## # A tibble: 10 x 2
##   carrier      n
##   <chr>   <int>
## 1 MQ      25397
## 2 AA      22558
## 3 UA       1693
## 4 9E       1044
## 5 B6        830
## 6 US        699
## 7 FL        187
## 8 DL        110
## 9 F9         50
## 10 WN         38
```

Answer: A majority of planes without missing tailnumbers belong to the MQ and AA carriers.